

**VISVESVARAYA TECHNOLOGICAL
UNIVERSITY**

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT

on

Machine Learning (23CS6PCMAL)

Submitted by

Dipesh Sah (1BM22CS092)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

Feb-2025 to Jun-2025

**B.M.S. College of Engineering,
Bull Temple Road, Bangalore 560019**
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Dipesh Sah (1BM22CS092)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Rashmi H Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
---	--

Index

Sl. No.	Date	Experiment Title	Page No.
1	17-3-2025	Write a python program to import and export data using Pandas library functions	4
2	17-3-2025	Demonstrate various data pre-processing techniques for a given dataset	6
3	24-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	15
4	11-3-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	17
5	11-3-2025	Build Logistic Regression Model for a given dataset	21
6	7-4-2025	Build KNN Classification model for a given dataset.	23
7	7-4-2025	Build Support vector machine model for a given dataset	26
8	21-4-2025	Implement Random Forest ensemble method on a given dataset.	31
9	5-5-2025	Implement Boosting ensemble method on a given dataset.	33
10	5-5-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	36
11	29-4-2025	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	39

Github Link:

Program 1

Write a python program to import and export data using Pandas library functions

Screenshot:

Date : 17 March, 2025

End-to-End Machine Learning.

Step. 1: Download the dataset in CSV format i.e. for housing.csv.

The attributes of the dataset being :

- i) Longitude
- ii) Latitude
- iii) housing_medianage
- iv) total_rooms
- v) total_bedrooms
- vi) population
- vii) household
- viii) median_income
- ix) median_house_value
- x) ocean_proximity .

Step. 2: Take a quick look at the data structure housing.head()

- print first five rows of database housing.info
- print all information about the dataset using housing.describe().
- print all the information i.e. count, max, min, mean, std, etc of all columns.

Step. 3: Print the histogram of the dataset by using housing.hist(bins=50, figsize(20,15))

Step. 4: Create a test & train set.

Divide the dataset into train & test.

Train is used to train the model.

Test is used to test the model.

train_set, test_set, train_test_split(housing, test_size=0.2) & Display the train_set & test_set.

Step. 5: Discover & visualize the data to gain insights.

(i) Visualize geographic data.
housing.plot(kind='scatter', x='longitude', y='latitude')

(ii) Discover the correlation matrix.
corr_matrix = housing.corr()

(iii) Discover the scatter matrix plot by scatter_matrix(housing[attributes], figsize(12, 12))

Step. 6: Prepare the data for machine learning algorithm.

housing = strat_train_set.drop("median_house_val")
housing_labels = strat_train_set["median_house_val"].copy()
This will make it possible to easily fix various transformations and for vehicle transformations work test.

Step. 7: Data Cleaning

- (i) Get rid of the corresponding directions.
- (ii) Get rid of whole attributes.
- (iii) Set the values to some values (zero, mean, etc)

Step. 8: Select and Train a model by using from sklearn.linear_model.

import linear regression
lin-reg = linear regression()
lin-reg.fit(housing_prepared, housing_labels)

Code:

```
import pandas as pd
# ----- Step 1: Import Data from CSV -----
# If you already have a CSV file, replace 'data.csv' with the filename.
# For this demo, let's create a DataFrame and save it to CSV first.
data = {
    'Name': ['Alice', 'Bob', 'Charlie'],
    'Age': [24, 30, 22],
    'City': ['New York', 'Los Angeles', 'Chicago']
}
# Create DataFrame
df = pd.DataFrame(data)
# Exporting DataFrame to CSV
df.to_csv('people.csv', index=False)
print("Data exported to 'people.csv'.")
# ----- Step 2: Import Data from CSV -----
imported_df = pd.read_csv('people.csv')
print("\nData imported from 'people.csv':")
print(imported_df)
# ----- Step 3: Export Data to Excel -----
df.to_excel('people.xlsx', index=False)
print("\nData exported to 'people.xlsx'.")
# ----- Step 4: Import Data from Excel -----
imported_excel_df = pd.read_excel('people.xlsx')
print("\nData imported from 'people.xlsx':")
print(imported_excel_df)
```

OUTPUT:

```
Data exported to 'people.csv'.
```

```
Data imported from 'people.csv':
      Name   Age       City
0    Alice    24  New York
1     Bob    30  Los Angeles
2  Charlie    22     Chicago
```

```
Data exported to 'people.xlsx'.
```

```
Data imported from 'people.xlsx':
      Name   Age       City
0    Alice    24  New York
1     Bob    30  Los Angeles
2  Charlie    22     Chicago
```

Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot:

Date : 17 March, 2025

End-to-End Machine Learning.

Step. 1: Download the dataset in CSV format i.e. for housing.csv. The attributes of the dataset being :

- i) Longitude
- ii) Latitude
- iii) housing_medianage
- iv) total_rooms
- v) total_bedrooms
- vi) population
- vii) household
- viii) median_income
- ix) median_house_value
- x) ocean_proximity .

Step. 2). Take a quick look at the data structure housing.head().

- print first five rows of database housing.info
- print all information about the dataset using housing.describe().
- print all the information i.e. count, max, min, mean, std, etc of all columns.

Step. 3) Print the histogram of the dataset by using housing.hist(bins=50, figsize(20,15)).

Step. 4) Create a test & train set.

Divide the dataset into train & test.

Train is used to train the model.

Test is used to test the model.

train_set, test_set, train_test_split(housing, test_size=0.2) & Display the train_set & test_set.

Step. 5) Discover & visualize the data to gain insights.

- (i) Visualize geographic data.
housing.plot(kind='scatter', x='longitude', y='latitude')
- (ii) Discover the correlation matrix.
corr_matrix = housing.corr()
- (iii) Discover the scatter matrix plot by scatter_matrix(housing[attributes], figsize(12,5))

Step. 6) Prepare the data for machine learning algorithm.

housing = strat_train_set.drop("median_house_value", axis=1)

housing_labels = strat_train_set["median_house_value"].copy()

This will make it possible to easily fix various transformations and for vehicle transformations work test.

Step. 7) Data Cleaning

- (i) Get rid of the corresponding dissections.
- (ii) Get rid of whole attributes.
- (iii) Set the values to some values (zero, mean, etc)

Step. 8) Select and Train a model by using from sklearn.linear_model.

import linear regression

lin-reg = linear regression()

lin-reg.fit(housing_prepared, housing_labels)

Code:

```
from google.colab import files
diabetes=files.upload()

from google.colab import files
adult_income=files.upload()

df1=pd.read_csv("Dataset of Diabetes .csv")
df1.head()

df2=pd.read_csv("adult.csv")
df2.head()
```

```

df1.info()
df2.info()
df1.describe()
df2.describe()

missing_values1 = df1.isnull().sum()
print(missing_values1)
missing_values2 = df2.isnull().sum()
print(missing_values2)

df1['Gender'] = df1['Gender'].replace('f', 'F')
ordinal_encoder = OrdinalEncoder(categories=[["F", "M"]])
df1["Gender_Encoded"] =
ordinal_encoder.fit_transform(df1[["Gender"]]) onehot_encoder =
OneHotEncoder()
encoded_data =
onehot_encoder.fit_transform(df1[["CLASS"]]) encoded_array
= encoded_data.toarray()
encoded_df = pd.DataFrame(encoded_array,
columns=onehot_encoder.get_feature_names_out(["CLASS"])) df_encoded = pd.concat([df1,
encoded_df], axis=1)
df1 = pd.concat([df1, encoded_df], axis=1)
df1.drop("CLASS", axis=1, inplace=True)
df1.drop("Gender", axis=1, inplace=True)
print(df2.head())
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
df_copy2 = df2
ordinal_encoder = OrdinalEncoder(categories=[["Male", "Female"]])
df_copy2["Gender_Encoded"] =
ordinal_encoder.fit_transform(df_copy2[["gender"]])
print(df_copy2[["gender", "Gender_Encoded"]])

onehot_encoder = OneHotEncoder()
encoded_data =
onehot_encoder.fit_transform(df2[["occupation", "workclass", "education
", "marital status", "relationship", "race", "native-country", "income"]])
encoded_array = encoded_data.toarray()
encoded_df =
pd.DataFrame(encoded_array,
columns=onehot_encoder.get_feature_names_out(["occupation", "workclass", "education
", "marital status", "relationship", "race", "native-country", "income"]))
df_encoded = pd.concat([df_copy2, encoded_df], axis=1)

df_encoded.drop("gender", axis=1, inplace=True)
df_encoded.drop("occupation", axis=1, inplace=True)
df_encoded.drop("workclass", axis=1, inplace=True)
df_encoded.drop("education", axis=1, inplace=True)
df_encoded.drop("marital-status", axis=1, inplace=True)
df_encoded.drop("relationship", axis=1, inplace=True)
df_encoded.drop("race", axis=1, inplace=True)
df_encoded.drop("native-country", axis=1, inplace=True)
df_encoded.drop("income", axis=1, inplace=True)

```

```

print(df_encoded.head())
normalizer = MinMaxScaler()
df_encoded[["fnlwgt","educational-num","capital-gain","capital-loss","hours-per-week"]] =
normalizer.fit_transform(df_encoded[["fnlwgt","educational-num","capital-gain","capital-
loss","hours-per week"]])
])
df_encoded.head()
normalizer = MinMaxScaler()
df1[["No_Pation","AGE","Urea","Cr" , "HbA1c" ,
"Chol","TG","HDL","LDL","VLDL","BMI"]] =
normalizer.fit_transform(df1[["No_Pation","AGE","Urea","Cr" , "HbA1c" ,
"Chol","TG","HDL","LDL","VLDL","BMI"]])
df1.head()

import os
import tarfile
import urllib
import pandas as pd
DOWNLOAD_ROOT = "https://raw.githubusercontent.com/ageron/handson-ml2/master/"
HOUSING_PATH = os.path.join("datasets", "housing")
HOUSING_URL = DOWNLOAD_ROOT + "datasets/housing/housing.tgz"

def fetch_housing_data(housing_url=HOUSING_URL, housing_path=HOUSING_PATH):
    os.makedirs(housing_path, exist_ok=True)
    tgz_path = os.path.join(housing_path, "housing.tgz")
    urllib.request.urlretrieve(housing_url, tgz_path)
    housing_tgz = tarfile.open(tgz_path)
    housing_tgz.extractall(path=housing_path)
    housing_tgz.close()
fetch_housing_data()

def load_housing_data(housing_path=HOUSING_PATH):
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)
housing = load_housing_data()

housing.head()

```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY

```
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   longitude        20640 non-null   float64
 1   latitude         20640 non-null   float64
 2   housing_median_age 20640 non-null   float64
 3   total_rooms      20640 non-null   float64
 4   total_bedrooms   20433 non-null   float64
 5   population       20640 non-null   float64
 6   households       20640 non-null   float64
 7   median_income    20640 non-null   float64
 8   median_house_value 20640 non-null   float64
 9   ocean_proximity  20640 non-null   object  
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
housing['ocean_proximity'].value_counts()
```

```
count
```

ocean_proximity	count
<1H OCEAN	9136
INLAND	6551
NEAR OCEAN	2658
NEAR BAY	2290
ISLAND	5

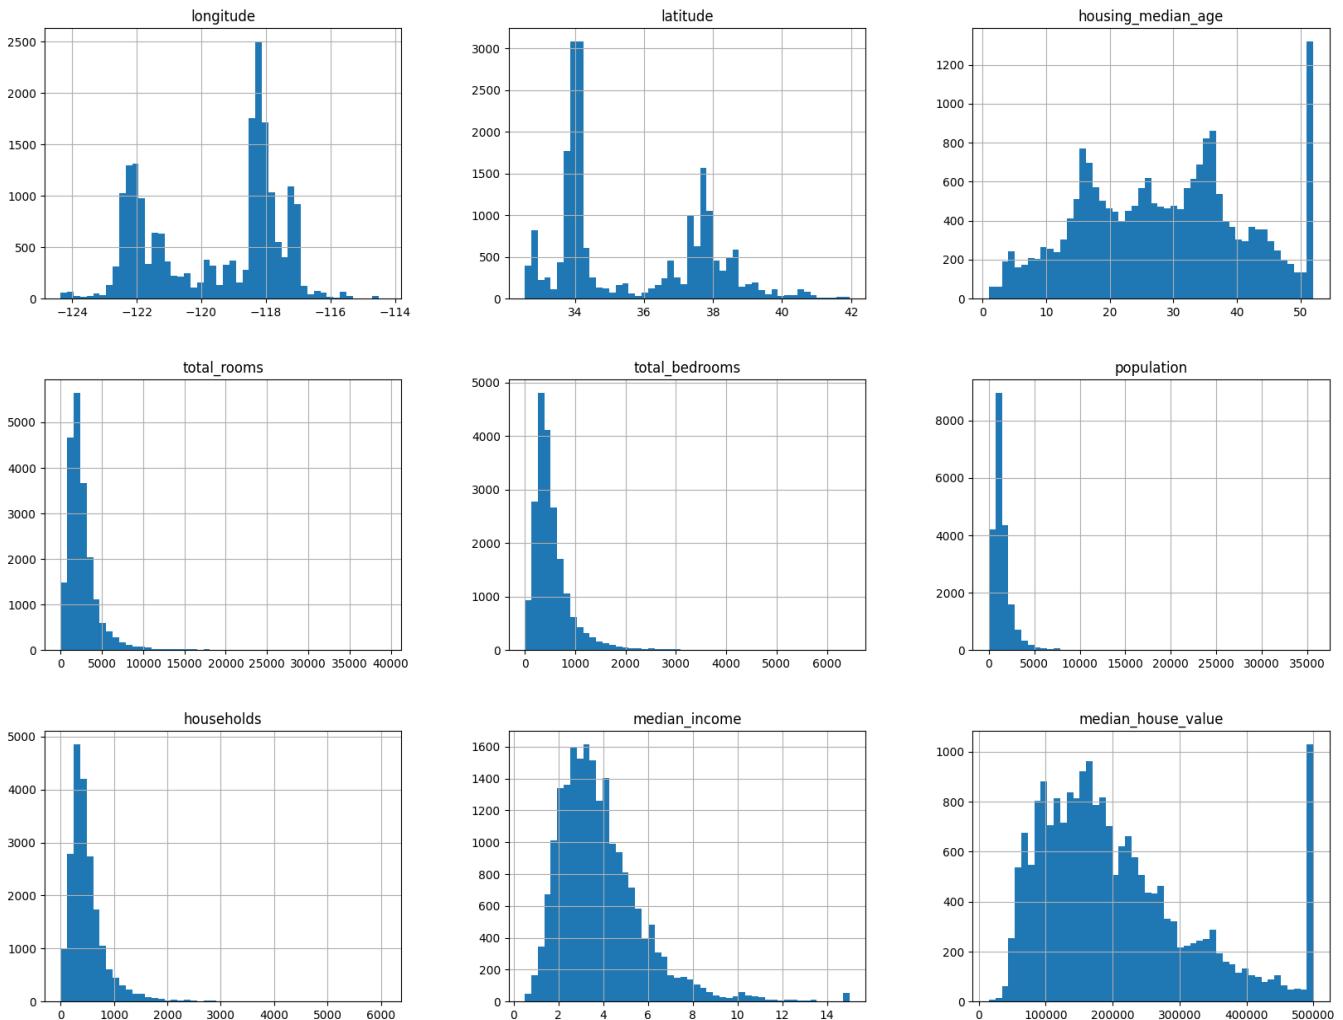
```
dtype: int64
```

```
housing.describe()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206855.816909
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115395.615874
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14999.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119600.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179700.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264725.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500001.000000

```
%matplotlib inline
```

```
import matplotlib.pyplot as plt
housing.hist(bins=50, figsize=(20,15))
plt.show()
```



```

import numpy as np
# def split_train_test(data, test_ratio):
#   np.random.seed(42)
#   shuffled_indices = np.random.permutation(len(data))
#   test_set_size = int(len(data) * test_ratio)
#   test_indices = shuffled_indices[:test_set_size]
#   train_indices = shuffled_indices[test_set_size:]
#   return data.iloc[train_indices], data.iloc[test_indices]

from sklearn.model_selection import train_test_split, StratifiedShuffleSplit

train_set, test_set = train_test_split(housing, test_size=0.2, random_state=42)
len(test_set)

```

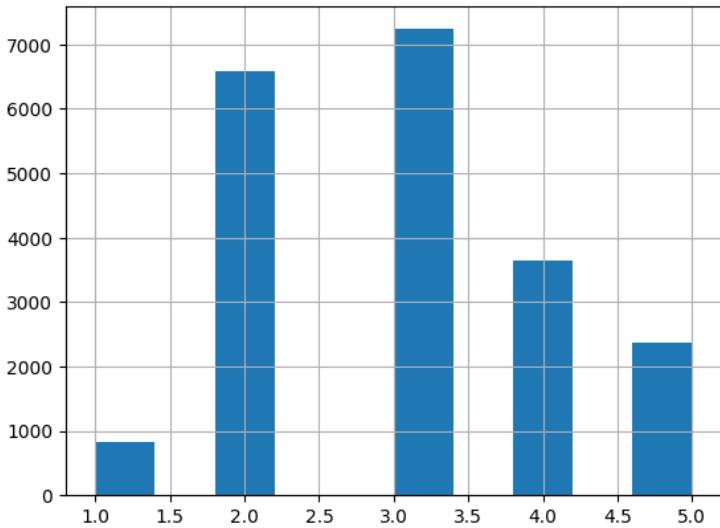
4128

```

housing['income_cat'] = pd.cut(housing['median_income'],
                               bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                               labels=[1, 2, 3, 4, 5])

```

```
housing['income_cat'].hist()
```

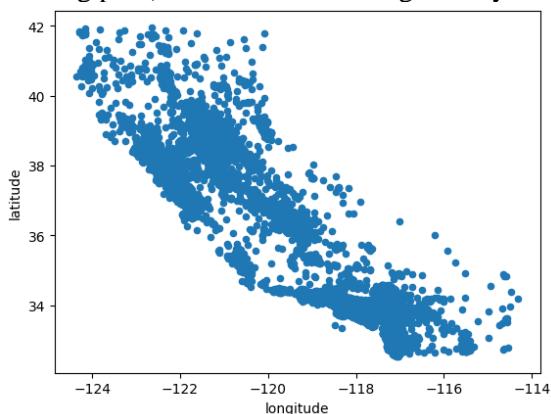


```
split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
for train_index, test_index in split.split(housing, housing['income_cat']):
    strat_train_set = housing.loc[train_index]
    strat_test_set = housing.loc[test_index]
```

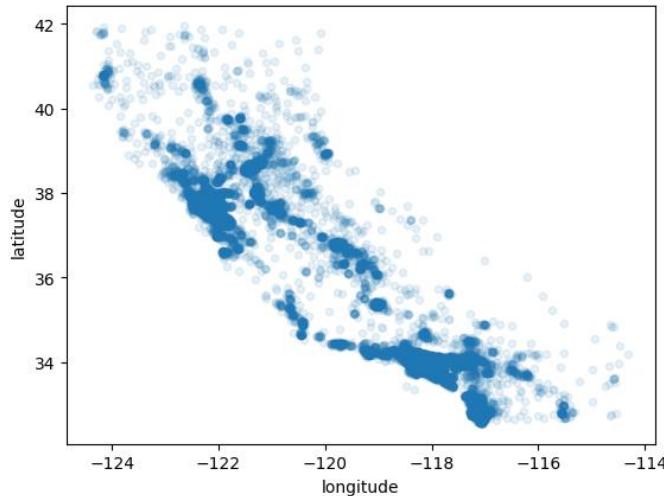
```
strat_test_set['income_cat'].value_counts() / len(strat_test_set)
```

income_cat	count
3	0.350533
2	0.318798
4	0.176357
5	0.114341
1	0.039971

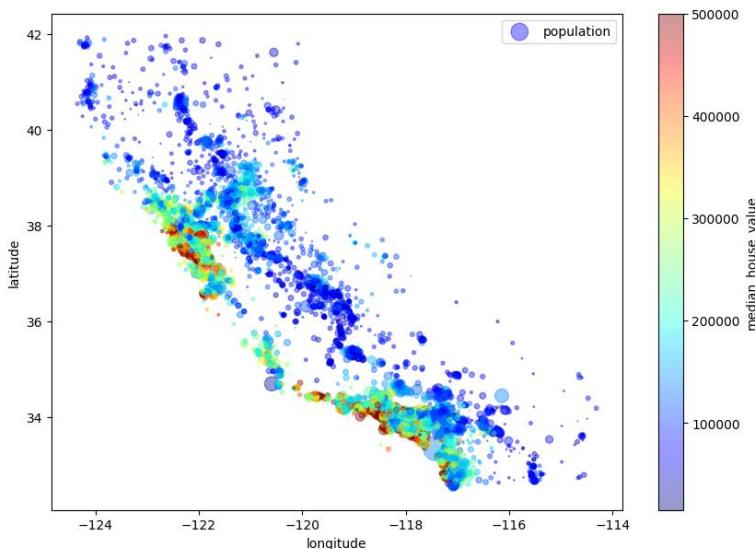
```
for set_ in (strat_train_set, strat_test_set):
    set_.drop('income_cat', axis=1, inplace=True)
housing = strat_train_set.copy()
housing.plot(kind='scatter', x='longitude', y='latitude')
```



```
housing.plot(kind='scatter', x='longitude', y='latitude', alpha=0.1)
```



```
housing.plot(kind='scatter', x='longitude', y='latitude', alpha=0.4,
             s=housing['population']/100, label='population', figsize=(10, 7),
             c='median_house_value', cmap=plt.get_cmap('jet'), colorbar=True)
plt.legend()
```



```
# Select only numerical features for correlation calculation
```

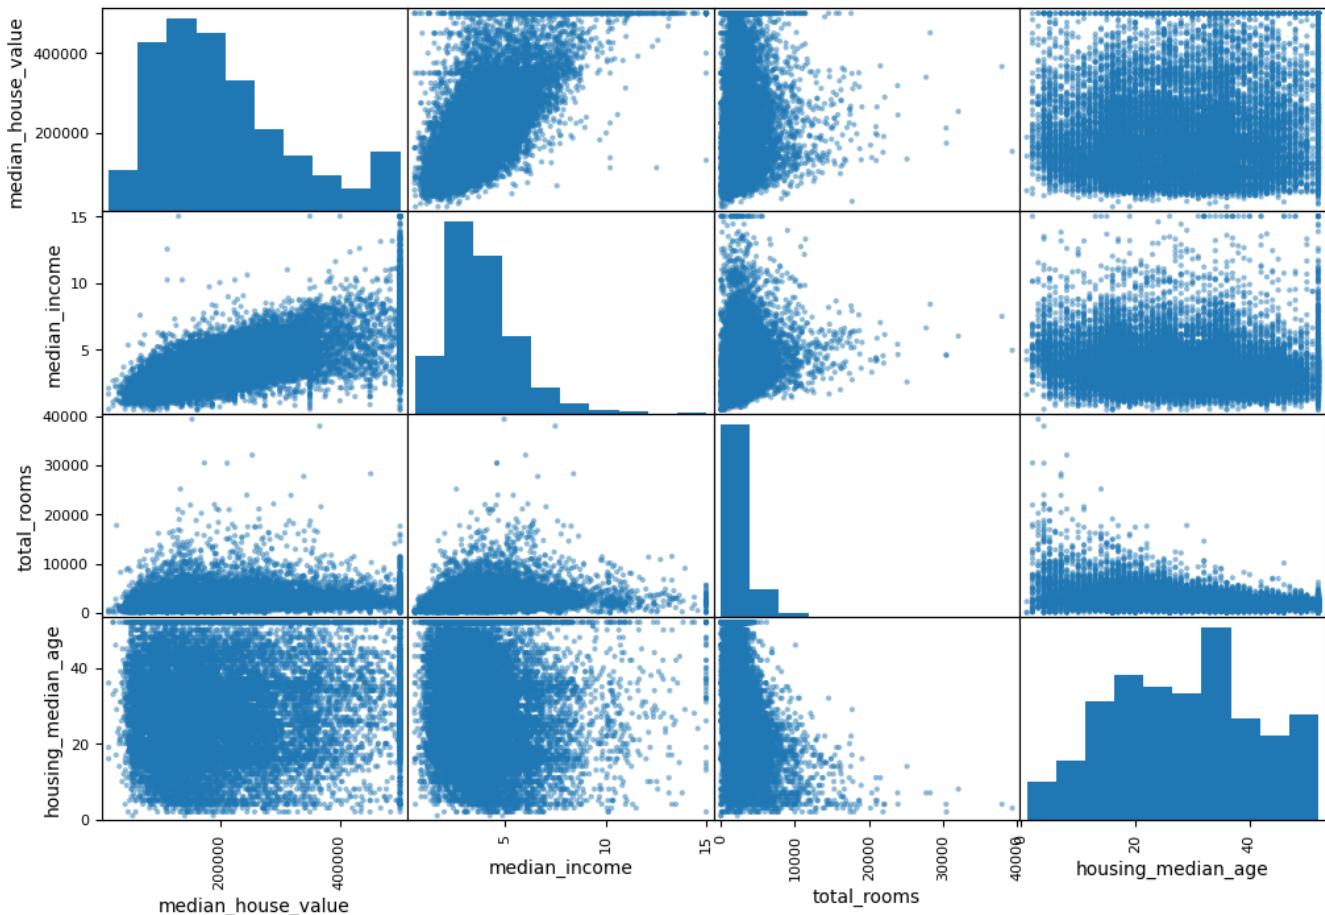
```
numerical_features = housing.select_dtypes(include=['number']) # Exclude non-numeric columns
corr_matrix = numerical_features.corr() # Calculate correlation for numerical features only
corr_matrix['median_house_value'].sort_values(ascending=False)
```

median_house_value	
median_house_value	1.000000
median_income	0.687151
total_rooms	0.135140
housing_median_age	0.114146
households	0.064590
total_bedrooms	0.047781
population	-0.026882
longitude	-0.047466
latitude	-0.142673

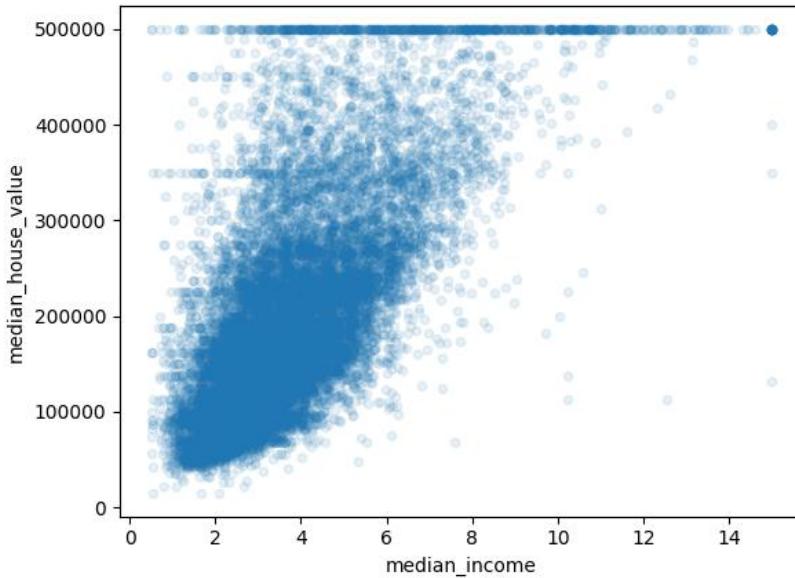
dtype: float64

```
from pandas.plotting import scatter_matrix
attributes = ['median_house_value', 'median_income', 'total_rooms', 'housing_median_age']
```

```
scatter_matrix(housing[attributes], figsize=(12, 8))
```



```
housing.plot(kind='scatter', x='median_income', y='median_house_value', alpha=0.1)
```



```
housing['rooms_per_household'] = housing['total_rooms']/housing['households']
housing['bedrooms_per_room'] = housing['total_bedrooms']/housing['total_rooms']
housing['population_per_household'] = housing['population']/housing['households']
```

```
# Select only numerical features for correlation calculation
numerical_features = housing.select_dtypes(include=['number']) # Exclude non-numeric columns
corr_matrix = numerical_features.corr() # Calculate correlation for numerical features only
corr_matrix['median_house_value'].sort_values(ascending=False)
```

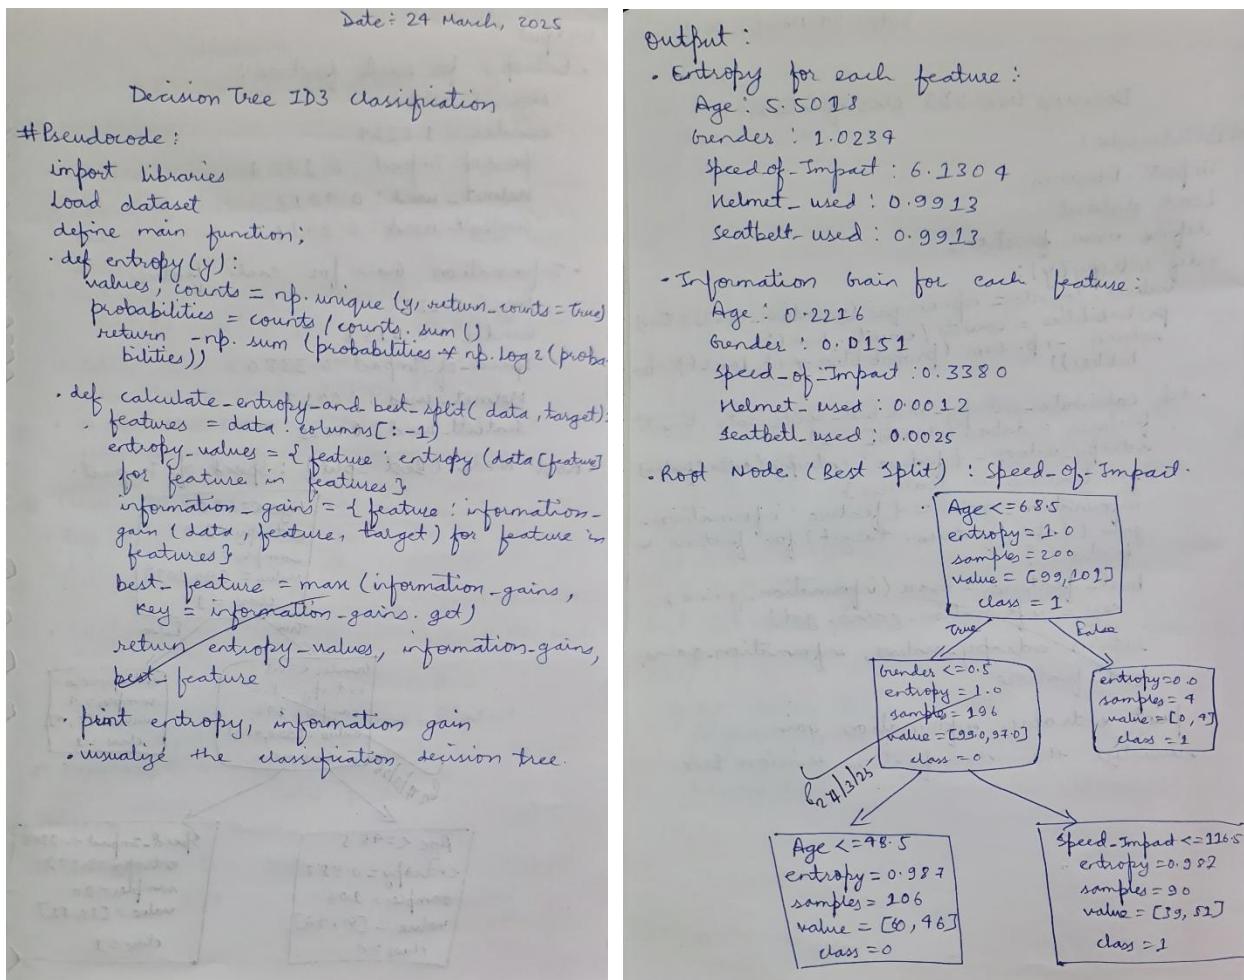
	median_house_value
median_house_value	1.000000
median_income	0.687151
rooms_per_household	0.146255
total_rooms	0.135140
housing_median_age	0.114146
households	0.064590
total_bedrooms	0.047781
population_per_household	-0.021991
population	-0.026882
longitude	-0.047466
latitude	-0.142673
bedrooms_per_room	-0.259952

dtype: float64

Program 3

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot:



Code:

```

import pandas as pd
import numpy as np
# Load the dataset
file_path = "/content/diabetes.csv"
df = pd.read_csv(file_path)
def entropy(df, col):
    ppos = (df[col] == 1).sum()
    pneg = len(df[col]) - ppos
    total = ppos + pneg
    if ppos == 0 or pneg == 0:
        return 0
    ppos /= total
    pneg /= total

```

```

    return -ppos * np.log2(ppos) - pneg * np.log2(pneg)
def information_gain(df, feature, target):
    total_entropy = entropy(df, target)
    values, counts = np.unique(df[feature], return_counts=True)
    weighted_entropy = sum((counts[i] / sum(counts)) * entropy(df[df[feature] == values[i]], target)
                           for i in range(len(values)))
    return total_entropy - weighted_entropy

def id3(df, features, target):
    # If all target values are the same, return that class
    if len(df[target].unique()) == 1:
        return df[target].iloc[0]

    # If no features left, return the most common target value
    if not features:
        return df[target].mode()[0]

    # Compute information gain for all features
    info_gains = {feature: information_gain(df, feature, target) for feature in features}

    # Print information gain for each feature
    print("Information Gain for each feature:")
    for feature, gain in info_gains.items():
        print(f"{feature}: {gain}")

    # Select the best feature
    best_feature = max(info_gains, key=info_gains.get)
    tree = {best_feature: {}}

    # Recur for each unique value in the selected feature
    for value in df[best_feature].unique():
        sub_df = df[df[best_feature] == value].drop(columns=[best_feature])
        tree[best_feature][value] = id3(sub_df, [f for f in features if f != best_feature], target)

    return tree

def print_tree(tree, indent=""):
    if not isinstance(tree, dict):
        print(indent + "→ " + str(tree))
        return
    for key, value in tree.items():
        print(indent + str(key))
        for sub_key in value:
            print(indent + " |--- " + str(sub_key))
            print_tree(value[sub_key], indent + " | ")

```

Compute the decision tree

```

features = [col for col in df.columns if col != "Outcome"]
decision_tree = id3(df, features, "Outcome")

print("Decision Tree:")
print_tree(decision_tree)

```

Program 4

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot

Date : 10 - March - 2025

Linear Regression

Linear Regression is a statistical method used to model the relationship between a dependent variable and one or more independent variables. Linear regression is also a type of supervised learning algorithm that learns from the labelled datasets and maps the data points with most optimized linear functions which can be used for prediction on new datasets.

It computes the linear relationship between the dependent variable and one or more independent features by fitting a linear equation with observed data.

linear equation with single variable:

$$\hat{Y} = \beta_0 + \beta_1 X + \epsilon_i$$

\hat{Y} → dependent
 X → independent
 β_0 → intercept
 β_1 → slope
 ϵ → error

① $\beta_1 = \frac{\bar{XY} - \bar{X}\bar{Y}}{\bar{X}^2 - \bar{X}^2}$ $\beta_0 = \bar{Y} - \beta_1 \bar{X}$
 using Mean Based formula

② $\beta_1 = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$
 using OLS method.

③ $\beta_1 = \frac{\sum n_i y_i - n \bar{y} \bar{x}}{\sum x_i^2 - n \bar{x}^2}$
 using std least square method

```
# compilation of working code for the same -
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

df = pd.read_csv("data_for_lr.csv")
df = df.iloc[:50]
X = df.iloc[:, 0].values
y = df.iloc[:, 1].values

# compute linear regression parameters using
# the least squares formula
X_mean, y_mean = np.mean(X), np.mean(y)

# Slope (m) & intercept (b)
m = np.sum((X - X_mean) * (y - y_mean)) / np.sum((X - X_mean) ** 2)
b = y_mean - m * X_mean

# Predicted values
y_pred = m * X + b

# Mean Squared Error
mse = np.mean((y - y_pred) ** 2)
print(f"Slope (m): {m}, Intercept (b): {b}, MSE: {mse}")

# Plot data points & regression line
plt.scatter(X, y, color='blue', label='Data Points')
plt.plot(X, y_pred, color='red', label='Regression Line')
plt.xlabel('X values')
plt.ylabel('Y values')
plt.legend()
plt.title("Linear Regression (Limited to 50 points)")
plt.show()
```

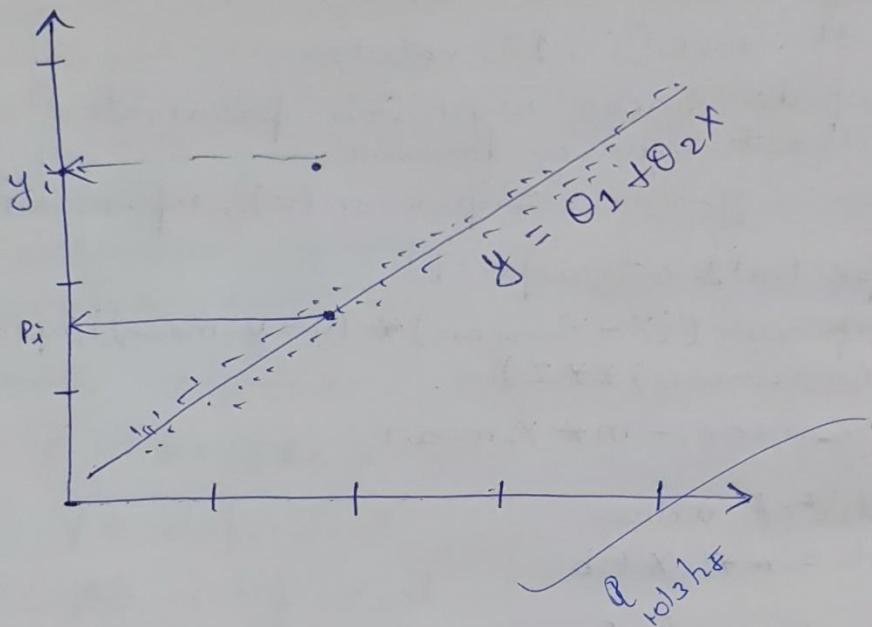
Output :-

Slope (m) : 0.992652844814438

Intercept (b) : -0.0738029467819

MSE : 6.4211389746497

Linear Regression
(Limited to 50 points)



Code:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
df = pd.read_csv("/content/sample_data/linear_regression_dataset.csv")
df1 = pd.read_csv("/content/sample_data/data_for_lr.csv")
X = df[['Experience (Years)']].values.flatten()
y = df['Salary ($1000s)'].values
# X = df1['x'].values.flatten()
# y = df1['y'].values.flatten()

X, y

(array([ 7,  4, 13, 15, 11,  8, 13,  5,  7, 10,  3,  7, 11, 11,  8,  5,  4,
       8,  8,  3,  6,  5,  2,  8, 12, 14,  6,  2, 12,  5]), array([20184, 10459, 40385, 43021, 33300, 22747, 39904, 16632, 19474,
       29082, 9558, 22753, 33047, 34547, 24747, 13975, 11806, 22189,
       25005, 9734, 19005, 13562, 5899, 25638, 35267, 42879, 17528,
       7202, 37556, 16890])))

mean_x = np.mean(X)
mean_y = np.mean(y)

numerator = np.mean(X * y) - mean_x * mean_y
denominator = np.mean(X**2) - mean_x**2

beta1 = numerator / denominator
beta0 = mean_y - beta1 * mean_x

y_pred = beta0 + beta1 * X

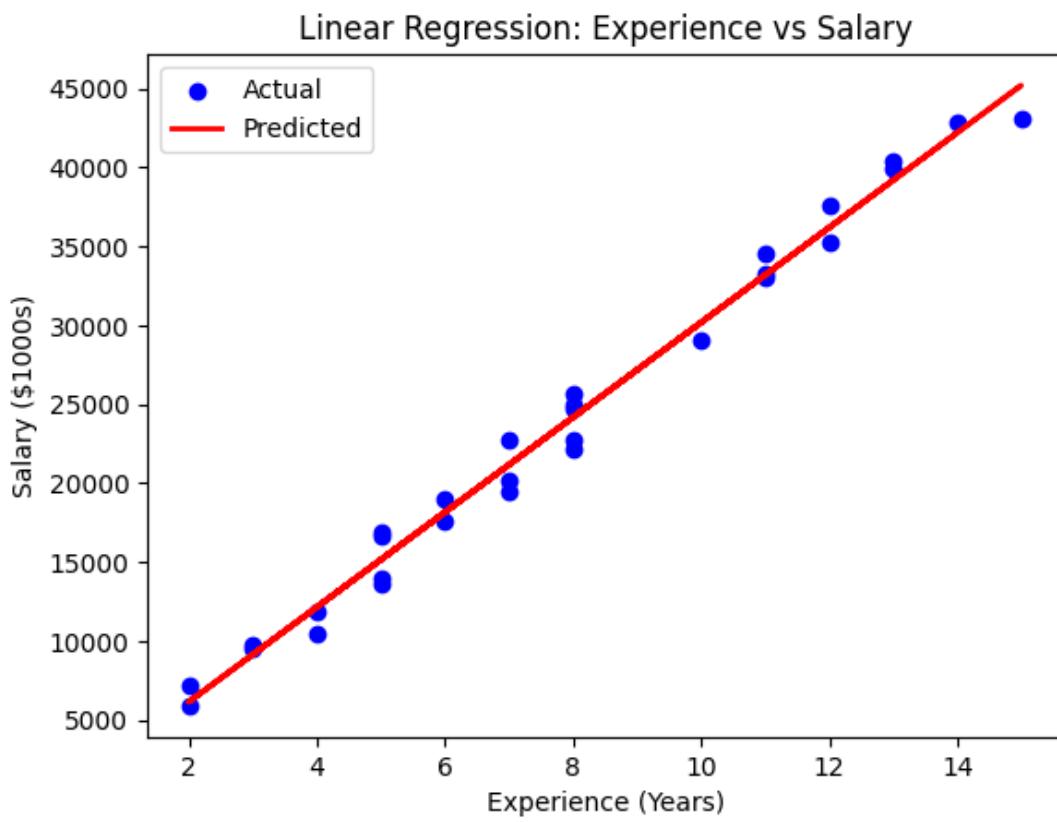
error = np.sum((y - y_pred)**2)
mean_x, mean_y

(7.766666666666667, 23465.833333333332)

print(f"Intercept (beta0): {beta0:.2f}")
print(f"Slope (beta1): {beta1:.2f}")
print(f"Error: {error:.2f}")

Intercept (beta0): 138.77
Slope (beta1): 3003.49
Error: 43353835.99

plt.scatter(X, y, color='blue', label='Actual')
plt.plot(X, y_pred, color='red', linewidth=2, label='Predicted')
plt.xlabel('Experience (Years)')
plt.ylabel('Salary ($1000s)')
plt.title('Linear Regression: Experience vs Salary')
plt.legend()
plt.show()
```



```
result = f"Linear Equation: Y = {beta0:.2f} + {beta1:.2f}X + {error:.2f}"  
result
```

```
'Linear Equation: Y = 138.77 + 3003.49X + 43353835.99'
```

Program 5

Build Logistic Regression Model for a given dataset

Code:

```
from google.colab import files
hr=files.upload()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder,
OneHotEncoder from sklearn.preprocessing import
StandardScaler, MinMaxScaler
from scipy import stats
from sklearn import linear_model
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

df1=pd.read_csv("HR_comma_sep.csv")
df1.head()
df1.isnull().sum()
plt.figure(figsize=(12, 6))
sns.barplot(x='Department', y='left', data=df1)
plt.title('Employee Retention Rate by Department')
plt.xlabel('Department')
plt.ylabel('Proportion of Employees Left')
plt.xticks(rotation=45, ha='right')
plt.show()

ohe = OneHotEncoder(handle_unknown='ignore',
sparse_output=False) department_encoded =
ohe.fit_transform(df1[['Department']])
department_encoded_df = pd.DataFrame(department_encoded,
columns=ohe.get_feature_names_out(['Department']))
df1 = pd.concat([df1, department_encoded_df], axis=1)
df1 = df1.drop('Department', axis=1)
ordinal_encoder = OrdinalEncoder(categories=[['low', 'medium', 'high']],
dtype=np.int64) salary_encoded =
ordinal_encoder.fit_transform(df1[['salary']])
df1['salary_encoded'] = salary_encoded
df1 = df1.drop('salary', axis=1)
df1.head()

correlation_matrix = df1.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
```

```

fmt=".2f") plt.title('Correlation Matrix of Features')
plt.show()
plt.figure(figsize=(8, 6))
sns.barplot(x='salary_encoded', y='left', data=df1)
plt.title('Impact of Employee Salary on Retention')
plt.xlabel('Salary Level (Encoded)')
plt.ylabel('Proportion of Employees Left')
plt.show()

df_copy = df1[['number_project', 'average_monthly_hours', 'time_spend_company',
'left','salary_encoded', 'satisfaction_level','Work_accident']]
df_copy.head()
X = df_copy.drop('left', axis=1)
y = df_copy['left']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the Logistic Regression model: {accuracy}")

from google.colab import files
zodata=files.upload()
zootype=files.upload()

zoo_data = pd.read_csv('zoo-data.csv')
zoo_class = pd.read_csv('zoo-class-type.csv')
merged_data      = pd.merge(zoo_data,      zoo_class,      left_on='class_type',
right_on='Class_Number') merged_data = merged_data.drop(['Animal_Names',
'Number_Of_Animal_Species_In_Class',
'Class_Number','class_type','animal_name'], axis=1)
X = merged_data.drop('Class_Type', axis=1)
y = merged_data['Class_Type']
print(merged_data.head())
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=np.unique(y_test)) disp.plot(cmap="Blues", values_format="d")
plt.title("Confusion Matrix")
plt.show()

```

Program 6

Build KNN Classification model for a given dataset.

Screenshot

Date: 07-04-2025

Lab - 4
* K-Nearest Neighbour

Algorithm:

- Choose k (number of neighbours to consider)
- Calculate distance (commonly Euclidean) between the new data point and all training points.
- Distance Metric (Euclidean):

$$d(p, q) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \dots + (p_n - q_n)^2}$$
- Select the k-nearest neighbours (those with the smallest distance).
- Assign the class based on majority vote among the k-neighbours.

How to choose k-value?

- Try different values for k (only choose odd no. to avoid ties)
- $k = \sqrt{n}$ $n = \text{no. of rows/ collection.}$
- Compute the accuracy / error rate for each value on a validation set.
- plot k vs accuracy / error rate.
- Choose the k with the highest accuracy / lowest error.

Accuracy = $\frac{\text{correct predictions}}{\text{Total predictions}}$

Error = 1 - Accuracy.

Dataset : Iris Dataset

Rows : 150 , columns = 6

Support Vector Machine (SVM)

⇒ Algorithm:

- Minimize $\frac{1}{2} \|w\|^2$
 - Subject to $y_i(w \cdot x_i + b) \geq 1$
- w = weights , b = bias
 x_i = input vector
 y_i = class label (+1/-1)

⇒ Feature Scaling

SVM is distance-based, it uses dot products.

It is the process of normalizing or standardizing the range of independent variables so that no single feature dominate others.

Standardization (Z-score)

$$n' = \frac{n - \mu}{\sigma}$$

Min-Max scaling

$$n' = \frac{n - n_{\min}}{n_{\max} - n_{\min}}$$

Code:

```
import csv
import math
import random
from collections import Counter
```

```

# ----- KNN Implementation -----

def load_dataset(filename):
    with open(filename, 'r') as f:
        data = list(csv.reader(f))
        header = data[0]
        rows = data[1:]
        for row in rows:
            for i in range(4):
                row[i] = float(row[i])
    return rows

def compute_mean_std(dataset):
    means, stds = [], []
    for i in range(4):
        col = [row[i] for row in dataset]
        mean = sum(col) / len(col)
        std = (sum((x - mean) ** 2 for x in col) / len(col)) ** 0.5
        means.append(mean)
        stds.append(std)
    return means, stds

def normalize_dataset(dataset, means, stds):
    for row in dataset:
        for i in range(4):
            row[i] = (row[i] - means[i]) / stds[i]
    return dataset

def split_dataset(dataset, test_size=0.2):
    random.shuffle(dataset)
    split_index = int(len(dataset) * (1 - test_size))
    return dataset[:split_index], dataset[split_index:]

def euclidean_distance(row1, row2):
    return math.sqrt(sum((row1[i] - row2[i]) ** 2 for i in range(4)))

def knn_predict(train, test_row, k):
    distances = []
    for train_row in train:
        dist = euclidean_distance(test_row, train_row)
        distances.append((train_row, dist))
    distances.sort(key=lambda x: x[1])
    neighbors = distances[:k]
    labels = [neighbor[0][4] for neighbor in neighbors]
    prediction = Counter(labels).most_common(1)[0][0]
    return prediction

def evaluate_model(train, test, k):
    correct = 0
    predictions = []
    for row in test:
        prediction = knn_predict(train, row, k)

```

```

predictions.append(prediction)
if prediction == row[4]:
    correct += 1
accuracy = correct / len(test)
return accuracy, predictions

filename = '/content/sample_data/IRIS.csv'
dataset = load_dataset(filename)

means, stds = compute_mean_std(dataset)

dataset = normalize_dataset(dataset, means, stds)

train_data, test_data = split_dataset(dataset)

k = 3
accuracy, predictions = evaluate_model(train_data, test_data, k)

print(f"Model Accuracy: {accuracy:.2f}\n\n")

print("\n--- Predict Iris Species from Your Input ---")
try:
    user_input = []
    user_input.append(float(input("Enter Sepal Length (cm): ")))
    user_input.append(float(input("Enter Sepal Width (cm): ")))
    user_input.append(float(input("Enter Petal Length (cm): ")))
    user_input.append(float(input("Enter Petal Width (cm): ")))

    for i in range(4):
        user_input[i] = (user_input[i] - means[i]) / stds[i]

    predicted_species = knn_predict(train_data, user_input, k)
    print(f"\n Predicted Iris Species: {predicted_species}")

except ValueError:
    print("Invalid input. Please enter numeric values for all measurements.")

Model Accuracy: 0.97

```

```

--- Predict Iris Species from Your Input ---
Enter Sepal Length (cm): 5.1
Enter Sepal Width (cm): 3.5
Enter Petal Length (cm): 1.4
Enter Petal Width (cm): 0.3

```

```
Predicted Iris Species: Iris-setosa
```

Program 7

Build Support vector machine model for a given dataset

Screenshot

Support Vector Machine (SVM)

⇒ Algorithm:

- Minimize $\frac{1}{2} \|w\|^2$
- Subject to $y_i(w \cdot x_i + b) \geq 1$

w = weights, b = bias
x_i = input vector
y_i = class label (+1/-1)

⇒ Feature scaling

SVM is distance-based, it uses dot products.

It is the process of normalizing or standardizing the range of independent variables so that no single feature dominate others.

Standardization (Z-score)

$x' = \frac{x - \mu}{\sigma}$ *6/4/2025*

Min-Max scaling

$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$

Code:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
def PlotData(x,y):
    pos = np.argwhere(y == 1)
    neg = np.argwhere(y == 0)

    plt.plot(x[pos, 0], x[pos, 1], linestyle="", marker='+', color='k')
    plt.plot(x[neg, 0], x[neg, 1], linestyle="", marker='o', color='y')
    plt.xlabel('Exam 1 score')
    plt.ylabel('Exam 2 score')
    plt.legend(['Admitted', 'Not admitted'], loc='upper right', numpoints=1)
    plt.figure()

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from sklearn.svm import SVC

df = pd.read_csv("/content/sample_data/IRIS.csv")

X = df[["petal_length", "petal_width"]].values
y = df["species"].values

label_map = {"Iris-setosa": 0, "Iris-versicolor": 1, "Iris-virginica": 2}
label_map_rev = {v: k for k, v in label_map.items()}
y_numeric = np.array([label_map_rev[label] for label in y])

svm_model = SVC(kernel="linear")
svm_model.fit(X, y_numeric)

x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 300),
                      np.linspace(y_min, y_max, 300))

Z = svm_model.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

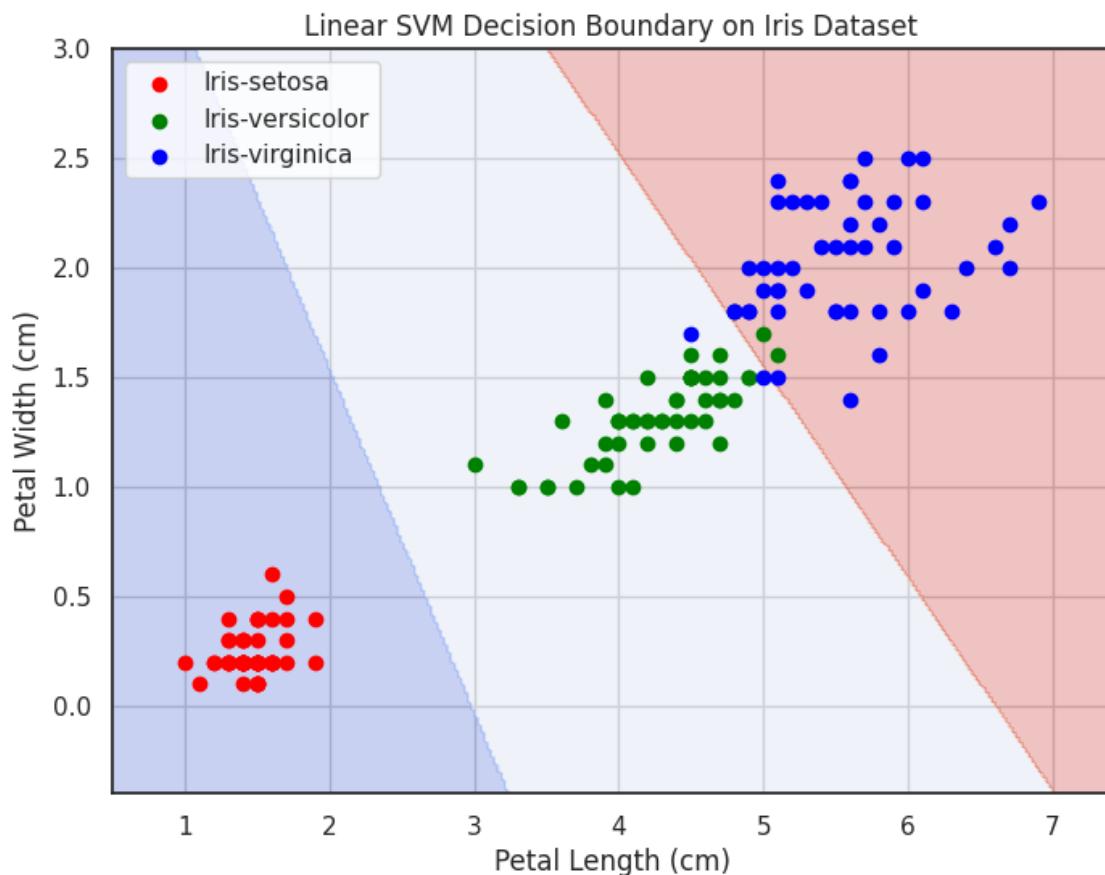
colors = ["red", "green", "blue"]

plt.figure(figsize=(8, 6))
plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.coolwarm)

for label in np.unique(y_numeric):
    plt.scatter(X[y_numeric == label, 0], X[y_numeric == label, 1],
                label=label_map_rev[label], color=colors[label])

plt.xlabel("Petal Length (cm)")
plt.ylabel("Petal Width (cm)")
plt.title("Linear SVM Decision Boundary on Iris Dataset")
```

```
plt.legend()  
plt.grid(True)  
plt.show()
```



```

import csv
import math
import random
from collections import Counter
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

def load_dataset(filename):
    df = pd.read_csv("/content/sample_data/IRIS.csv")
    X = df[["petal_length", "petal_width"]].values
    y = df["species"].values

    label_map = {"Iris-setosa": 0, "Iris-versicolor": 1, "Iris-virginica": 2}
    label_map_rev = {v: k for k, v in label_map.items()}
    y_numeric = np.array([label_map[label] for label in y])

class SimpleLinearModel:
    def __init__(self, learning_rate=0.01, n_iters=1000):
        self.lr = learning_rate
        self.n_iters = n_iters
        self.activation = self._unit_step_func
        self.weights = None
        self.bias = None

    def fit(self, X, y):
        n_samples, n_features = X.shape
        self.weights = np.zeros(n_features)
        self.bias = 0

        y_ = np.array([1 if i > 0 else 0 for i in y])

        for _ in range(self.n_iters):
            for idx, x_i in enumerate(X):
                linear_output = np.dot(x_i, self.weights) + self.bias
                y_predicted = self.activation(linear_output)

                update = self.lr * (y_[idx] - y_predicted)
                self.weights += update * x_i
                self.bias += update

    def predict(self, X):
        linear_output = np.dot(X, self.weights) + self.bias
        y_predicted = self.activation(linear_output)
        return y_predicted

    def _unit_step_func(self, x):
        return np.where(x>=0, 1, 0)

```

```

svm_model = SimpleLinearModel()
svm_model.fit(X, y_numeric)

x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 300),
                      np.linspace(y_min, y_max, 300))

Z = svm_model.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

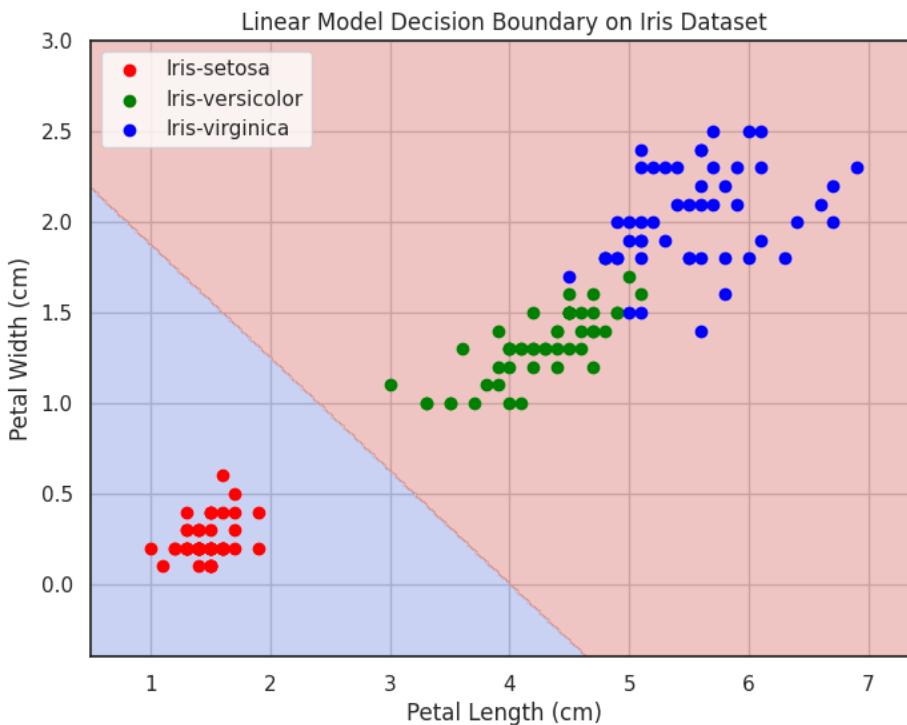
colors = ["red", "green", "blue"]

plt.figure(figsize=(8, 6))
plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.coolwarm)

for label in np.unique(y_numeric):
    plt.scatter(X[y_numeric == label, 0], X[y_numeric == label, 1],
                label=label_map_rev[label], color=colors[label])

plt.xlabel("Petal Length (cm)")
plt.ylabel("Petal Width (cm)")
plt.title("Linear Model Decision Boundary on Iris Dataset")
plt.legend()
plt.grid(True)
plt.show()

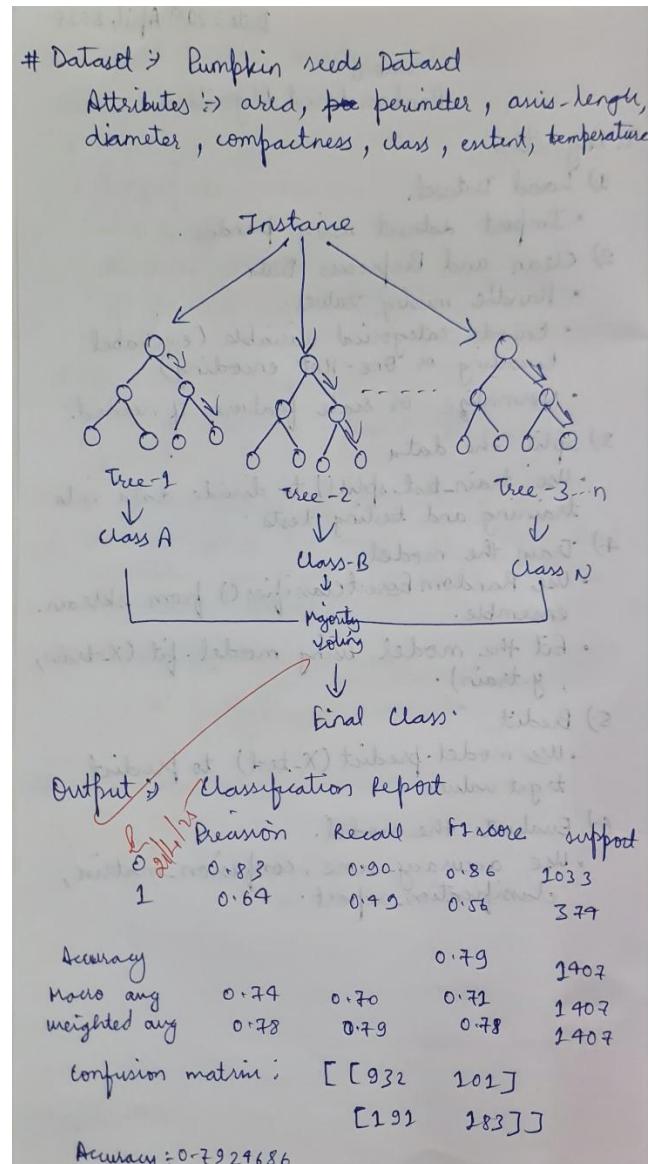
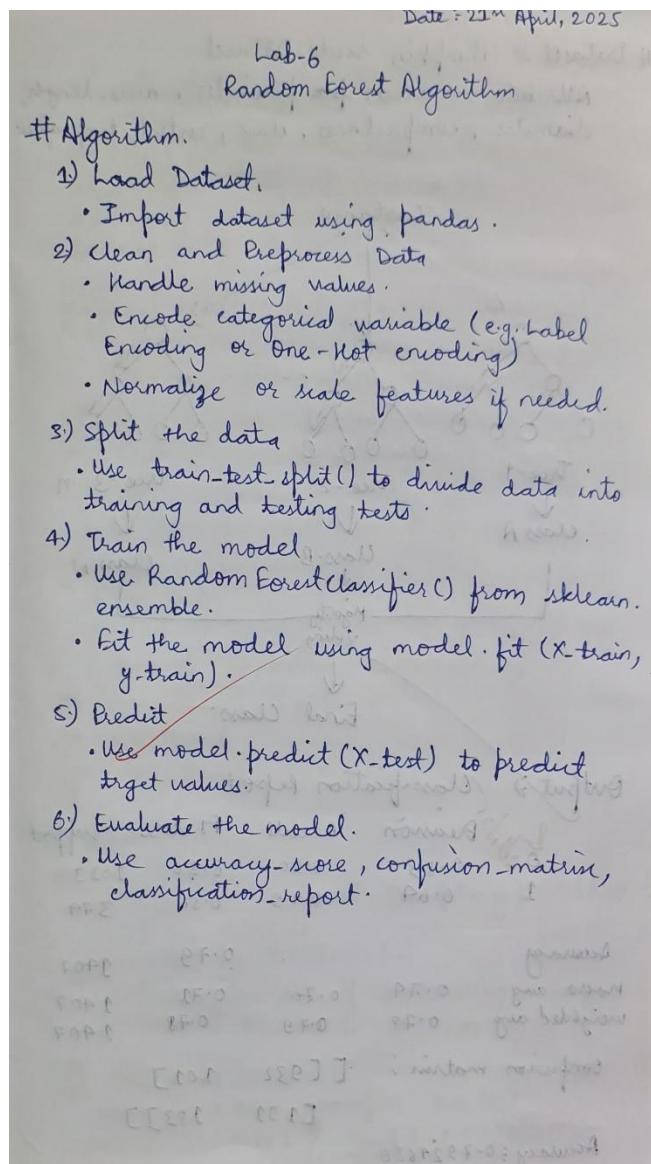
```



Program 8

Implement Random Forest ensemble method on a given dataset.

Screenshot



Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset
```

```

file_path = "/content/Pumpkin_Seeds_Dataset.xlsx"
df = pd.read_excel(file_path, sheet_name='Pumpkin_Seeds_Dataset')

# Separate features and target
X = df.drop(columns=['Class'])
y = df['Class']

# Encode the target labels
le = LabelEncoder()
y_encoded = le.fit_transform(y)

# Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y_encoded, test_size=0.2, random_state=42, stratify=y_encoded
)

# Train the Random Forest classifier
rf_model = RandomForestClassifier(n_estimators=100, max_depth=None, random_state=42)
rf_model.fit(X_train, y_train)

# Predict on test data
y_pred = rf_model.predict(X_test)

# Evaluate the model
print(" ✅ Accuracy:", accuracy_score(y_test, y_pred))
print("\n 📈 Classification Report:")
print(classification_report(y_test, y_pred, target_names=le.classes_))

# Plot the confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=le.classes_, yticklabels=le.classes_)
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.tight_layout()
plt.show()

# Plot top 10 feature importances
importances = rf_model.feature_importances_
features = X.columns
indices = importances.argsort()[:-1][:-10] # Top 10

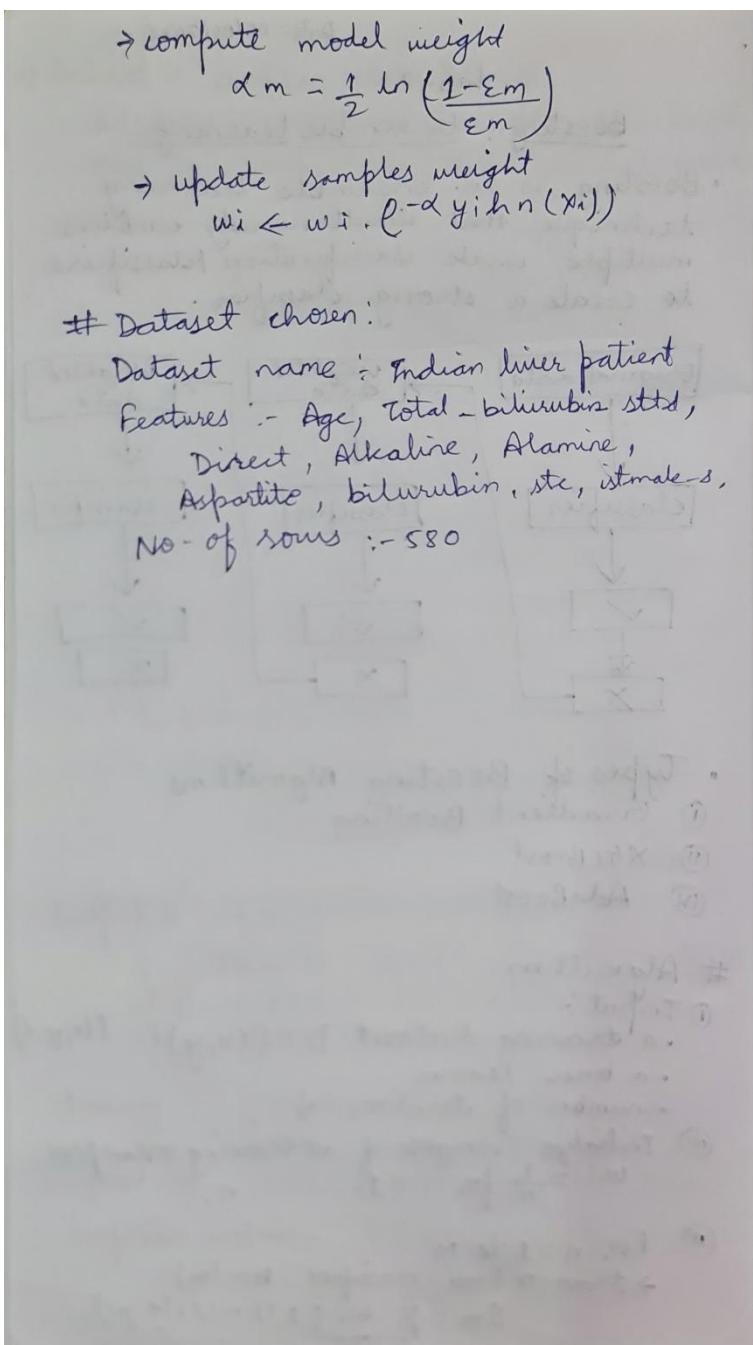
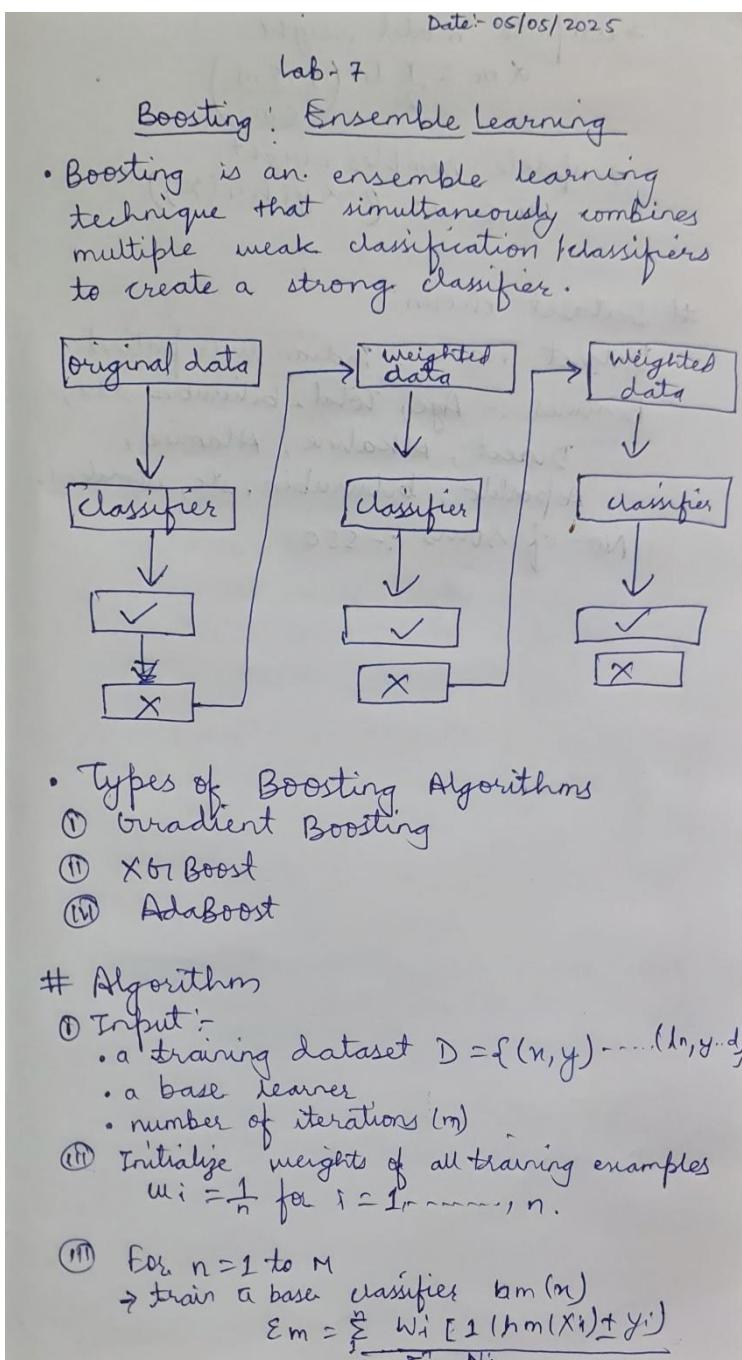
plt.figure(figsize=(10, 6))
sns.barplot(x=importances[indices], y=features[indices], palette="viridis")
plt.title("Top 10 Feature Importances")
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()

```

Program 9

Implement Boosting ensemble method on a given dataset.

Screenshot



Code:

```
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score

from sklearn.tree import DecisionTreeRegressor
from sklearn.tree import DecisionTreeClassifier

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier as KNN

from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import mean_squared_error as MSE

from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import VotingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingRegressor

SEED = 1
# Dataset
liver = pd.read_csv('/content/sample_data/indian_liver_patient_preprocessed.csv', index_col = 0)
X = liver.drop('Liver_disease', axis = 1)
y = liver['Liver_disease']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=SEED)

# Split data into 80% train and 20% test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
stratify=y,
random_state=SEED)

liver.head()

# Import AdaBoostClassifier
from sklearn.ensemble import AdaBoostClassifier

# Instantiate dt
dt = DecisionTreeClassifier(max_depth=2, random_state=1)

# Instantiate ada
ada = AdaBoostClassifier(estimator=dt, n_estimators=180, random_state=1)
```

```
# Fit ada to the training set
ada.fit(X_train, y_train)

# Compute the probabilities of obtaining the positive class
y_pred_proba = ada.predict_proba(X_test)[:,1]

# Import roc_auc_score
#from sklearn.metrics import roc_auc_score

# Evaluate test-set roc_auc_score
ada_roc_auc = roc_auc_score(y_test, y_pred_proba)

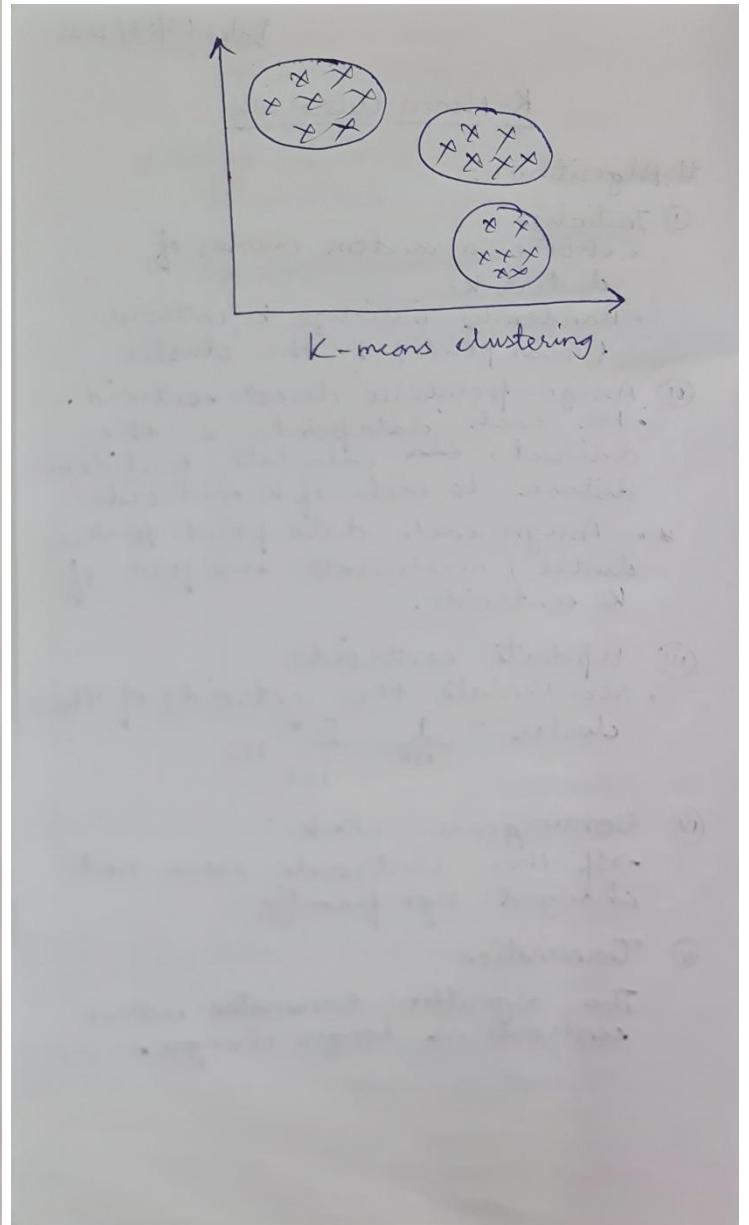
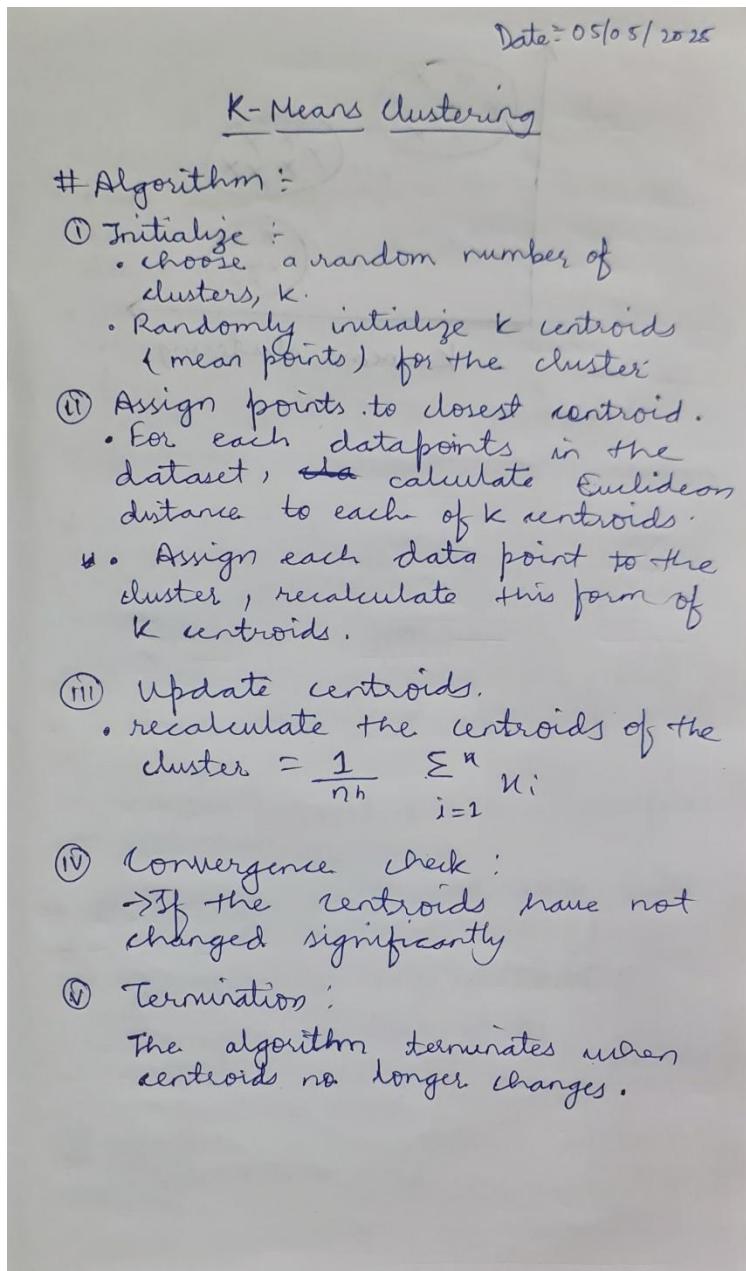
# Print roc_auc_score
print('ROC AUC score: {:.2f}'.format(ada_roc_auc))

ROC AUC score: 0.70
```

Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Screenshot



Code:

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # for data visualization
import seaborn as sns # for statistical data visualization
%matplotlib inline
```

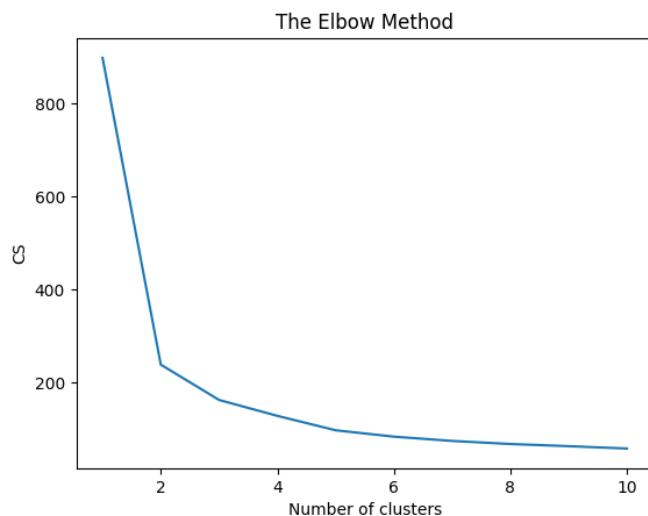
```

import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('/content/sample_data/Live.csv')
df.drop(['Column1', 'Column2', 'Column3', 'Column4'], axis=1, inplace=True)
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=2, random_state=0)
kmeans.fit(X)
kmeans.cluster_centers_
kmeans.inertia_
labels = kmeans.labels_
# check how many of the samples were correctly labeled
correct_labels = sum(y == labels)

print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
Result: 4288 out of 7050 samples were correctly labeled.
print('Accuracy score: {:.2f}'.format(correct_labels/float(y.size)))
Accuracy score: 0.61
from sklearn.cluster import KMeans
cs = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter = 300, n_init = 10, random_state = 0)
    kmeans.fit(X)
    cs.append(kmeans.inertia_)
plt.plot(range(1, 11), cs)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('CS')
plt.show()

```



```
from sklearn.cluster import KMeans
```

```

kmeans = KMeans(n_clusters=2,random_state=0)

kmeans.fit(X)

labels = kmeans.labels_

# check how many of the samples were correctly labeled

correct_labels = sum(y == labels)

print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))

print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))
Result: 4288 out of 7050 samples were correctly labeled.
Accuracy score: 0.61

kmeans = KMeans(n_clusters=3, random_state=0)

kmeans.fit(X)

# check how many of the samples were correctly labeled
labels = kmeans.labels_

correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))
Result: 4066 out of 7050 samples were correctly labeled.
Accuracy score: 0.58



---


kmeans = KMeans(n_clusters=4, random_state=0)

kmeans.fit(X)

# check how many of the samples were correctly labeled
labels = kmeans.labels_

correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." % (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.format(correct_labels/float(y.size)))
Result: 4112 out of 7050 samples were correctly labeled.
Accuracy score: 0.58

```

Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA) method.

Screenshot

Date = 12/05/2025

Principle Component Analysis (PCA)

- PCA is an unsupervised machine learning technique used for dimensionality reduction.
- It transforms the original features into a new set of features called principle components, which are linear combinations of the original features.
- Uses of PCA
 - Reduces noise & redundancy.
 - Improves performance of Machine learning algorithms.
 - helps in visualization, plotting high dimensional data in 2D or 3D..
- Algorithm:
Input : Dataset X with n samples and m features.
Output : Transformed data with reduced dimensions ($k \leq m$)

① Standardize the dataset :
For each feature in X :
subtract the mean and divide by standard deviation.

② compute the covariance matrix:
 $\text{cov Matrix} = \text{covariance_matrix}(X)$

③ compute eigen values & eigen vectors:
 $[\text{eigen values}, \text{eigen vectors}] = \text{eigen decomposition}(\text{cov matrix})$

④ sort eigen vectors by descending eigen values, sort eigen values and select top k :
select corresponding top k eigen vectors (principle components)

⑤ transform the original data :
 $\text{Reduced Data} = X * \text{principle components}$
return Reduced Data

Code:

```
from google.colab import files
heart=files.upload()
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from scipy import stats
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, OneHotEncoder from sklearn.model_selection import
train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score from sklearn.preprocessing
import StandardScaler
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.decomposition import PCA
df1=pd.read_csv("heart (1).csv")
df1.head()
text_cols = df1.select_dtypes(include=['object']).columns
label_encoder = LabelEncoder()
for col in text_cols:
    df1[col] =
    label_encoder.fit_transform(df1[col])
print(df1.head())
X = df1.drop('HeartDisease', axis=1)
y = df1['HeartDisease']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) scaler = StandardScaler()
X_train =
scaler.fit_transform(X_train) X_test =
scaler.transform(X_test)
# Support Vector Machine
svm_model = SVC(kernel='linear', random_state=42)
svm_model.fit(X_train, y_train)
svm_predictions = svm_model.predict(X_test)
svm_accuracy = accuracy_score(y_test, svm_predictions)
print(f"SVM Accuracy: {svm_accuracy}")
# Logistic Regression
lr_model = LogisticRegression(random_state=42)
lr_model.fit(X_train, y_train) lr_predictions = lr_model.predict(X_test) lr_accuracy = accuracy_score(y_test,
lr_predictions)
print(f"Logistic Regression Accuracy: {lr_accuracy}")
# Random Forest
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)
rf_accuracy = accuracy_score(y_test, rf_predictions)
print(f"Random Forest Accuracy: {rf_accuracy}")
```

```

models = {
    "SVM": svm_accuracy,
    "Logistic Regression":
        lr_accuracy, "Random Forest":
            rf_accuracy
    }
best_model = max(models, key=models.get)
print(f"\nBest Model: {best_model} with accuracy {models[best_model]}")
pca = PCA(n_components=0.95)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)
svm_model_pca = SVC(kernel='linear', random_state=42)
svm_model_pca.fit(X_train_pca, y_train)
svm_predictions_pca = svm_model_pca.predict(X_test_pca)
svm_accuracy_pca = accuracy_score(y_test, svm_predictions_pca)
print(f"SVM Accuracy (with PCA): {svm_accuracy_pca}")
lr_model_pca = LogisticRegression(random_state=42)
lr_model_pca.fit(X_train_pca, y_train)
lr_predictions_pca = lr_model_pca.predict(X_test_pca)
lr_accuracy_pca = accuracy_score(y_test, lr_predictions_pca)
print(f"Logistic Regression Accuracy (with PCA): {lr_accuracy_pca}")
rf_model_pca = RandomForestClassifier(random_state=42)
rf_model_pca.fit(X_train_pca, y_train)
rf_predictions_pca = rf_model_pca.predict(X_test_pca)
rf_accuracy_pca = accuracy_score(y_test, rf_predictions_pca)
print(f"Random Forest Accuracy (with PCA): {rf_accuracy_pca}")
models_pca = {
    "SVM": svm_accuracy_pca,
    "Logistic Regression": lr_accuracy_pca,
    "Random Forest": rf_accuracy_pca
}
best_model_pca = max(models_pca, key=models_pca.get)
print(f"\nBest Model (with PCA): {best_model_pca} with accuracy {models_pca[best_model_pca]}")

```